

Building predictors for the game mia

Falko Benezan

ID: 3617296

falko.benezan@
student.uni-tuebingen.de

Alexander Diegel

ID: 3980486

alexander.diegel@
student.uni-tuebingen.de

Abstract

This term paper deals with the topic of predicting different artificial intelligence approaches based on the game mia. It starts with an short introduction to artificial intelligence. Then, the description of the game mia on which the project is based on. Afterwards, the strategies of the different implemented artificial intelligences (AI) are explained. It follows an explanation of our approach predicting the behavior of the different AIs. Finally, the results of the experiments are discussed.

1 Introduction

Artificial intelligence has become an increasingly important field in computer science and other areas such as automotive (self-driving cars) or security (face-detection). In computer games artificial intelligence reaches new stages of success (Google bot AlphaGo for the game go). In this term paper we now want to answer the question if it is possible to predict different artificial intelligence approaches based on the game mia. The main achievement would be the accuracy of correct predictions if a player (AI) is willing to lie or is telling the correct value of the dice. The question if the actual value is lesser or greater than the called one seems not as important as the prediction of a lie itself.

2 The Game

Mia is a simple dice game that is played with two dices and a flat bottomed container (or a dice cup). At the beginning each player has a certain amount of lives (e.g. five). The first player rolls the dices but keeps their values hidden from the other players. He then can decide if he wants to tell the truth to the next player and announce a value that was

actually rolled. Alternatively he can lie and announce a greater or lesser value than the rolled one. But each player has to announce a greater value than the previous player. The next player (who still has not seen the actual values) can now believe the passer, call the passer a liar and look on the dice or pass the dice to the next player (still without looking) announcing a higher value. A player loses a life if he called the previous one a liar and looked on the values to find out that they are what the previous player has announced or even higher. Otherwise the previous player loses a life. The higher value of the roll is multiplied by ten and then added to the other die (a 4 and a 2 is 42). The **scoring** is from highest to lowest: 21 (Mia), 66, 55, 44, 33, 22, 11, 65, 64, 63, 62, 61, 54, 53, 52, 51, 43, 42, 41, 32, 31. If a player announces mia the next player either believes him, give up (without looking at the dices) and loses one life. Or he may look at the dice. If it was actually mia then he loses two lives if it was not, the previous player loses a life. (For further information see (mia, 2016).)

3 Implementation of the Game

For the implementation we use C++ and Qt to build the GUI-Application, as IDE we used Qt-Creator. To implement an artificial intelligence according to the support vector machine learner we used the OpenCV library. The ingame AI is used to distinguish between lies and the truth of announced values. For the prediction we used the last announced value, the actual value and an indicator to show the start of a new game. To train the SVM we additionally used an indicator if the actual (value that previous player announced) value was a lie.

The difference between the strategy for liar detection and gameplay follows in section 4.

4 Setup and Strategies

This section will describe the setup and different strategies used for this study. First of all there are different strategies (or player behaviors/ player types) that play the game and create game data. Each of those strategies has its own behavior in deciding whether the previous player lies at them or tells the truth. Moreover, each strategy has its own behavior in announcing a new value. Some will always tell the truth if this is possible, others will try to lie permanently. In the end we want to show the performance of this strategies meaning the liar detection of the player behaviors itself (how good can they predict if a previous player lied at them). Then we train an out of game predictor/classifier (see section 5) in python based on the game data to show if there is a way to beat the implemented strategies/player behaviors.

To measure the performance of the python classifier, different strategies or player behaviors were used in the form of some artificial intelligence approaches. For data acquisition we used a homogeneous set of players in our game implementation and let it generate game data. Therefore the output of one AI is the input of the following AI in the game. For each turn we reported:

- player number
- previous player number
- if the turn is the first one in the round (new game)
- new announced value
- actual rolled value
- previous announced value
- does the player look onto the dices or not

To train the SVM classifier we used values a real player would get plus the information if the announcement was the truth or a lie. Saying we used the actual announced value and the previous announced value of each turn in the game together with the label if this announcement was true or false (called *lierlabel*). When it comes to predict unseen test data we will give the SVM classifier only the previous and actual announced dice values (this is what a player in the real world would get). The SVM is then able to tell if the current announcement was the truth or a lie.

Each dataset is split up to training data and test data. So we train the python classifier for the different approaches by using cross validation (determination of parameter C) on the training data and measure the performance (generalization) on the test data. The test set is 30 percent of the data set, the others are training data. In the end we will compare the results of the implemented AIs detecting lies by taking the indicator that the ingame AI (player behavior) will look at the values (saying the former player was a liar) and the predictions of our python classifier (see section 6). The detection of lies made by the ingame AI (player behaviors) is called *sdetect* in table 2. The prediction (made by the python classifier) if an ingame AI is lying or not with its current announcement is called *pdetect* in table 2.

We examined three different types of artificial intelligence approaches/players behaviors. First the statistical approach where the AI acts very straight forward. Then an AI with a certain degree of randomness in its call and look behavior. At the end we implemented a learning AI with different calling behaviors to examine if the python classifier is able to learn the strategy of a learning player. Last we wanted to know how many samples it takes until there is convergence of accuracy in the python classifier.

4.1 Statistic Approach

Table 1 shows basic statistic probabilities of the game mia. Note: if the value (call/announcement of previous player) is 21 then the current player can either look up the dices and possibly loses two lives or he can save himself by rolling a 21 as well.

Value	31	32	41	42	43	51	52
%	94	89	83	78	72	67	61
Value	53	54	61	62	63	64	65
%	56	50	44	39	33	28	22
Value	11	22	33	44	55	66	21
%	19	17	14	11	8	6	6

Table 1: Shows the probability of getting a higher dice result than the announced value.

The artificial intelligence/behavior based on the statistical approach is quite easy. It calls the previous player a liar if the probability of lying is greater than the probability to say the truth. This AI always says the truth if possible, otherwise

it takes a random value greater than the last announced value when it rolls the dices.

4.2 Approach with certain degree of randomness (Primitive Random)

This AI will call with a certain probability a value that is greater than the actual rolled value if the value is lower than the one of the previous player. Also the AI will look on the dices with a probability that arises with decreasing beat probability (see table 1) of a announced value (a 31 will not be looked at with probability near 100 percent, 21 is looked at with 100 percent). This AI always says the truth if it is possible (meaning the rolled value is greater than the last announced value). Otherwise this AI will call a random value.

4.3 A SVM learning approach

The next step was to create a new type of artificial intelligence that changes its behavior during the history/time of the game. Therefore, we used a SVM implementation with different behaviors by calling a value (meaning lie or tell the truth). Each SVM implementation will be retrained after each new sample. This will give us a rather dynamic behavior in looking at the dices for all versions of this SVM strategy. The calling behavior for the different versions are listed below:

4.3.1 SVM variant 1

This variant of the AI will announce the true dice values or uses the next possible value greater than the last announced one.

4.3.2 SVM variant 2

This AI variant will announce the true dice values if it is possible, else it calls a random value greater than the last announced one.

4.3.3 SVM variant 3

This AI will always call a possible random value that is independent from the actual rolled dice value.

4.3.4 SVM variant 4

This AI will always announce a lower or equal value than the actual rolled value if it is possible. Otherwise it uses a random value greater than the last announced one. This AI is willing to lie whenever it is possible.

5 The Classifier

To predict the different strategical approaches we used a support vector machine (SVM) classifier with radial basis functions. A linear classifier would not have been sufficient due to the complexity of data, but with the SVM-classifier adequate results could be expected. By using SVM we are interested in separating two (or more) classes by a separating hyperplane with maximal margin. The margin is defined with respect to the training points as the minimal distance between the hyperplane and a training point. (See (Schoellkopf and Smola, 2002, 187–227).)

6 Discussion of Results

An overview of the maintained results can be seen in table 2. In all but one case the trained classifier was able to predict the liar labels better than the AIs. Moreover we can conclude that the classifier was able to learn the strategy of the different approaches quite well.

strategy	sdetect	C	pdetect
statistic	0.473	0.0001	0.863
primitive	0.585	0.5	0.849
SVM1	0.964	2.0	0.971
SVM2	0.810	0.6	0.797
SVM3	0.844	3.0	0.971
SVM4	0.642	2.0	0.989

Table 2: Obtained results in the experiments. Strategy, strategies liar detection (sdetect), parameter C for classifier, classifier's liar detection (pdetect) on test data. The classifier is a python SVM implementation.

The artificial intelligence based on the statistic approach (see section 4.1) has the worst results of all our experiments. Our classifier was able to achieve a result of 86% and more or less learn the AIs strategy. The approach with a certain degree of randomness (see section 4.2) was able to detect around 60% (the classifier again beat the ingame AI and was able to learn their strategy). The learning of the SVM helps to improve the ingame predictions and learn the strategy of the other players (the ingame results are quite as good as the our classifier implemented in python). Summing up we can say that the ingame AIs results are the better the less randomness exist in the calling behavior of the other players. The classifier was able to learn nearly all strategies good. The following convergences consider the convergence of the

python classifier not the ingame AI convergence.

For the statistical approach convergence takes place around 4000 samples if we allow one percent of deviation in both directions (see fig. 7). In fig. 6 the convergence can be seen more easily. It is again about 4000 samples. Afterwards there is only little variation of the accuracy. For the SVM approaches can be seen that variant 1 and variant 3 converge to the same value of accuracy (0.97). Convergence for both variants is about 2000 samples. Variant 2 of the SVM approach has a lower level of accuracy meaning that this one is harder to predict correctly than the others (under the given data). The convergence of variant 2 is at around 3000 samples. Variant 4 has the highest rate of accuracy with 98 percent and converges about 2500 samples (compare fig. 5).

Comparing the convergence of the ingame AI (fig. 8) and the python classifier (fig. 1) we can say that both converge near to 1500 samples. The strategy of a liar has been detected by both ingame and python classifier.

Figures 1 to 4 show the classification of the ingame artificial intelligence based on the SVM learner. The green color represents values that are going to be accepted as the truth by other players. Blue areas represent that the actual value is detected as a lie. On the horizontal axis there is the actual value. On the vertical axis there is the last value¹. On the left half the running game is represented and on the right side new game predictions are depicted. The origin of the image is on the top left side.




Figure 1: Classification of SVM-AI variant 1.

Figure 1 shows that the SVM-AI (see section 4.3.1) can learn the calling strategy of saying the next greater value. The blue points show that the next greater value is quite often a simple lie.

¹The axis are in numerical order. Therefore, non valid dice values for the mia game are represented.




Figure 2: Classification of SVM-AI variant 2

Figure 2 has a green area between the blue dots. That is because the values around 31 and 42 are often true and the AI accepts this calls (see section 4.3.2).




Figure 3: Classification of SVM-AI variant 1

In fig. 3 you can see clearly that the variance of the liars classification increases due to the random calling procedure of the third variant (see section 4.3.3). And since the AI is likely to lie there are blue points on the right side of the image, meaning that start values are considered to be likely a lie.




Figure 4: Classification of SVM-AI variant 1

In fig. 4 where the calling behavior is to say a smaller value (see section 4.3.4) if possible the variance of the liar calling is a bit smaller than in section 4.3.3.

7 Supplemental Material

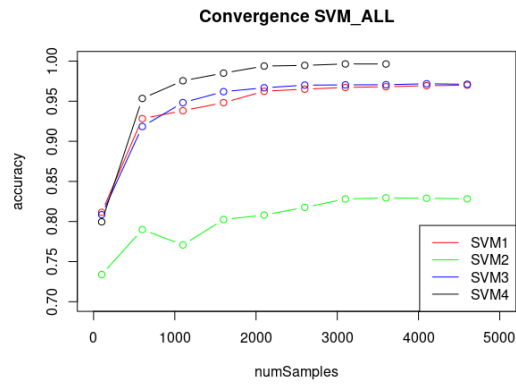


Figure 5: Comparison of predictions based on all four SVM variants data sets (player behaviors) of the python classifier.

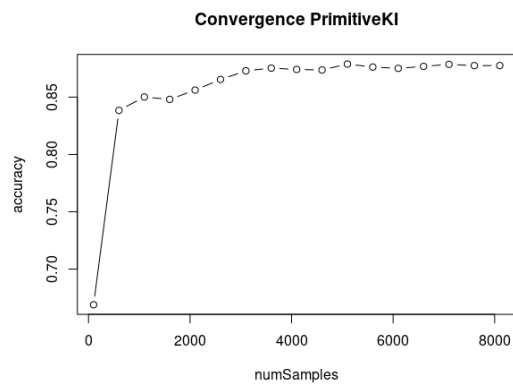


Figure 6: Convergence of the python classifier on data generated by the approach with certain degree of randomness.




Figure 7: Convergence of python classifier on data generated by the statistical approach.

Convergence of ingame SVM1




Figure 8: Convergence of the ingame SVM in variant 1.

References

- [mia2016] 2016. Mia (game). [https://en.wikipedia.org/wiki/Mia_\(game\)](https://en.wikipedia.org/wiki/Mia_(game)).
- [Schoellkopf and Smola2002] Bernhard Schoellkopf and Alexander J. Smola. 2002. *Learning with Kernels*. Massachusetts Institute of Technology, Cambridge, Massachusetts, London, England.